

# Regression 2


Masashi Sugiyama, Takashi Ishida

[sugi@k.u-tokyo.ac.jp](mailto:sugi@k.u-tokyo.ac.jp), [ishi@k.u-tokyo.ac.jp](mailto:ishi@k.u-tokyo.ac.jp)

<http://www.ms.k.u-tokyo.ac.jp>

# Regression = Function Approximation

2



Learn a function that is close to the underlying function w/ training samples

# Linear-in-parameter model

3

$$f_{\theta}(\mathbf{x}) = \sum_{j=1}^b \theta_j \phi_j(\mathbf{x})$$



$\{\phi_j(\mathbf{x})\}_{j=1}^b$   
: basis functions

## ■ Linear model:

$$f_{\theta}(\mathbf{x}) = \sum_{j=1}^n \theta_j K(\mathbf{x}, \mathbf{x}_j)$$

Gaussian kernel


$$K(\mathbf{x}, \mathbf{c}) = \exp \left( -\frac{\|\mathbf{x} - \mathbf{c}\|^2}{2h^2} \right)$$



# Least squares regression

$$\min_{\theta} \sum_{i=1}^n \left( f_{\theta}(x_i) - y_i \right)^2$$

- Minimize the squared error between training outputs:



# Regularization


5

$$\min_{\boldsymbol{\theta}} \left[ \underbrace{\frac{1}{2} \sum_{i=1}^n \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2}_{\text{goodness of fit for training output}} + \underbrace{\frac{\lambda}{2} \|\boldsymbol{\theta}\|^2}_{\text{penalty term for preventing param values from becoming too large (regularization)}} \right]$$

- The model strikes a good balance between fitting the training output well and keeping the parameter values small.
- Also called  $\ell_2$ -Regularized regression

# Cross validation

- Split training samples  $Z = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  into  $k$  groups:  $\{Z_i\}_{i=1}^k$
- Use samples from groups excluding  $Z_i$  and learn  $\theta$  (fix  $\lambda, h$ ).
- Use the remaining  $Z_i$  to check the test error.
- Repeat this for all  $i \in [k]$ , and return the mean of the test errors.





# Schedule

- |    |       |                  |     |       |                           |
|----|-------|------------------|-----|-------|---------------------------|
| 1. | 04/8  | Introduction     | 8.  | 06/17 | Deep learning 3           |
| 2. | 04/15 | Regression 1     | 9.  | 06/24 | Semi-supervised learning  |
| 3. | 04/22 | Regression 2     |     |       |                           |
| ●  | 04/30 | Cancelled        | 10. | 07/01 | Language models           |
| 4. | 05/13 | Classification 1 | 11. | 07/08 | Representation learning 1 |
| 5. | 05/20 | Classification 2 |     |       |                           |
| 6. | 05/27 | Deep learning 1  | 12. | 07/15 | Representation learning 2 |
| ●  | 06/03 | No lecture       |     |       |                           |
| 7. | 06/10 | Deep learning 2  | 13. | 07/22 | Advanced topics           |

# Sparsity of the model

- Having a large number of parameters makes computation more challenging.
- Example: with kernel models, it is computationally difficult when the number of training samples is large since #params = #training samples:

$$f_{\theta}(\mathbf{x}) = \sum_{j=1}^n \theta_j K(\mathbf{x}, \mathbf{x}_j)$$

- If many parameter values are set to zero, the computation becomes easier and more interpretable.



# Sparsity of the model


9

## ■ Naive approach 1:

- Decide to not use some of the parameters.
- When we have  $d$  params, we have  $2^d$  different ways of choosing params.
- Not a realistic approach when  $d$  is too large.

## ■ Naive approach 2:

- First use  $\ell_2$ -regularized least squares.
- Simply choose the params that have a small absolute value, and squash it to zero.
- May suffer from rounding errors.



# Contents

10

1. Sparse Regression
  1.  $\ell_1$ -constrained least squares regression
  2. Solving  $\ell_1$ -constrained LS
  3. Various extensions
2. Robust Regression


# $\ell_1$ -constrained least squares regression

- Constrain the model to be within  $\ell_1$  hyper-cube



$$\min_{\boldsymbol{\theta}} \frac{1}{2} \sum_{i=1}^n \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2 \quad \text{subject to } \|\boldsymbol{\theta}\|_1 \leq R$$

$$R \geq 0$$

$$\|\boldsymbol{\theta}\|_1 = \sum_{j=1}^b |\theta_j|$$



$\ell_2$ -constrained  
LS regression




$\ell_1$ -constrained  
LS regression


# Why do we get a sparse solution?

12

- The solution tends to be on one of the coordinate axes.



$\ell_2$ -constrained LS  
regression




$\ell_1$ -constrained LS  
regression

- Also called **sparse regression** or **LASSO** (Least Absolute Shrinkage and Selection Operator)

# Why do we get a sparse solution?

13

- This does not always happen! For example, when both variables are essentially important:



# Example


- Gaussian kernel model:

$$f_{\theta}(\mathbf{x}) = \sum_{j=1}^n \theta_j K(\mathbf{x}, \mathbf{x}_j)$$


Will explain  
algorithms soon!

$$K(\mathbf{x}, \mathbf{c}) = \exp\left(-\frac{\|\mathbf{x} - \mathbf{c}\|^2}{2h^2}\right)$$


Ordinary LS



$\ell_2$ -CLS



$\ell_1$ -CLS



- The results of  $\ell_1$ -CLS and  $\ell_2$ -CLS are similar
- Out of the 50 params in  $\ell_1$ -CLS, 38 params are zero!




# Feature Selection

- If we perform sparse learning for a **linear model w.r.t. the input**, then some input variables will be cancelled out.

$$f_{\theta}(x) = \theta^{\top} x \quad x = (x^{(1)}, \dots, x^{(d)})^{\top}$$

- Will be able to select the features useful for prediction automatically.
- **Example:** automatic selection of important genes
- Instead of considering all  $2^d$  combinations, we just need to decide  $\lambda$  in sparse learning.



# Contents

16

1. Sparse regression
  1.  $\ell_1$ -constrained least squares regression
  2. Solving  $\ell_1$ -constrained LS
  3. Various extensions
2. Robust regression


# Deriving the solution

- Equivalent expression with  $\ell_1$  hyper-cube constraint:

$$\hat{\boldsymbol{\theta}} = \operatorname{argmin}_{\boldsymbol{\theta}} \left[ \frac{1}{2} \sum_{i=1}^n \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2 + \lambda \sum_{j=1}^b |\theta_j| \right]$$


$\lambda (\geq 0)$ : constant decided by  $R$

- Instead of choosing  $R$ , we can specify  $\lambda$ .
- However, since the absolute value is not differentiable at the origin, the above optimization problem cannot be solved easily.





# Approximate gradient method

- $\theta \leftarrow \theta - \varepsilon \nabla J(\theta)$




- Approximate the derivative of the absolute value



- But it is unstable and does not work well in practice because many solutions are zero.

# Upper bound of $\ell_1$ -norm

- In order to consider a different direction, we first recall that the  $\ell_1$ -norm is:  $\|\boldsymbol{\theta}\|_1 = \sum_{j=1}^b |\theta_j|$
- Upper bound:  $|\theta_j| \leq \frac{\theta_j^2}{2c_j} + \frac{c_j}{2}$  for  $c_j > 0$
- Why? Apply inequality of arithmetic & geometric mean to RHS. Equality holds when  $\theta_j = \pm c_j$ .



- If we use the current param  $\tilde{\theta}_j \neq 0$  for  $c_j$ ,

$$|\theta_j| \leq \frac{\theta_j^2}{2|\tilde{\theta}_j|} + \frac{|\tilde{\theta}_j|}{2}$$

- If we regard  $|\theta_j| = 0$  for  $\tilde{\theta}_j = 0$ , then the upper bound for parameter  $\tilde{\theta}_j \neq 0$ :

$$|\theta_j| \leq \frac{|\tilde{\theta}_j|^\dagger}{2} \theta_j^2 + \frac{|\tilde{\theta}_j|}{2}$$

$|\theta|^\dagger = 1/|\theta|$ :  
generalized inverse

- Original objective:

$$J(\boldsymbol{\theta}) = J_{\text{LS}}(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1$$

- Upper bound of  $J(\boldsymbol{\theta})$ :

$$\tilde{J}(\boldsymbol{\theta}) = J_{\text{LS}}(\boldsymbol{\theta}) + \frac{\lambda}{2} \boldsymbol{\theta}^\top \tilde{\Theta}^\dagger \boldsymbol{\theta} + \lambda \sum_{j=1}^b |\tilde{\theta}_j|/2$$

$\tilde{\Theta}$ : diagonal matrix with  $|\tilde{\theta}_1|, \dots, |\tilde{\theta}_b|$  on diag       $\tilde{\Theta}^\dagger$ : diagonal matrix with  $|\tilde{\theta}_i|^\dagger$  on the diag



# Minimizing $\tilde{J}(\boldsymbol{\theta})$

$$\tilde{J}(\boldsymbol{\theta}) = J_{\text{LS}}(\boldsymbol{\theta}) + \frac{\lambda}{2} \boldsymbol{\theta}^\top \tilde{\boldsymbol{\Theta}}^\dagger \boldsymbol{\theta} + \text{Constant}$$

- If we are using a linear-in-parameter model  $f_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x})$  then:

$$J_{\text{LS}}(\boldsymbol{\theta}) = \frac{1}{2} \|\boldsymbol{\Phi} \boldsymbol{\theta} - \mathbf{y}\|^2.$$

- Recall from last week:  $\boldsymbol{\Phi} = [\boldsymbol{\phi}_j(\mathbf{x}_i)]_{j,i}$  is design matrix.
- The solution of minimizing  $\tilde{J}(\boldsymbol{\theta})$ :

$$\hat{\boldsymbol{\theta}} = (\boldsymbol{\Phi}^\top \boldsymbol{\Phi} + \lambda \tilde{\boldsymbol{\Theta}}^\dagger)^{-1} \boldsymbol{\Phi}^\top \mathbf{y}$$

- Note that we can ignore  $\sum_{j=1}^b |\tilde{\theta}_j|/2$  in the minimization problem since it is a constant w.r.t.  $\boldsymbol{\theta}$ .

# Algorithm

22

- 1. Initialize parameter  $\theta$ .
- 2. Form current  $\theta$ , derive  $\Theta$ :

$$\Theta \longleftarrow \text{diag}(|\theta_1|, \dots, |\theta_b|).$$

- 3. Update  $\theta$ :

$$\theta \longleftarrow (\Phi^T \Phi + \lambda \Theta^\dagger)^{-1} \Phi^T y.$$

- Repeat 2 and 3 until convergence.

This is called: **iteratively reweighted shrinkage**


# A closer look into the updates

23

- Some properties:
  - Since quadratic function is tangent to original function when  $\theta = \tilde{\theta}$ :  $J(\tilde{\theta}) = \tilde{J}(\tilde{\theta})$ .
  - Since  $\hat{\theta}$  is the minimizer of  $\tilde{J}$ :  $\tilde{J}(\tilde{\theta}) \geq \tilde{J}(\hat{\theta})$ .
  - Since  $\tilde{J}(\theta)$  is upper bound of  $J$ :  $\tilde{J}(\hat{\theta}) \geq J(\hat{\theta})$ .

- To summarize:

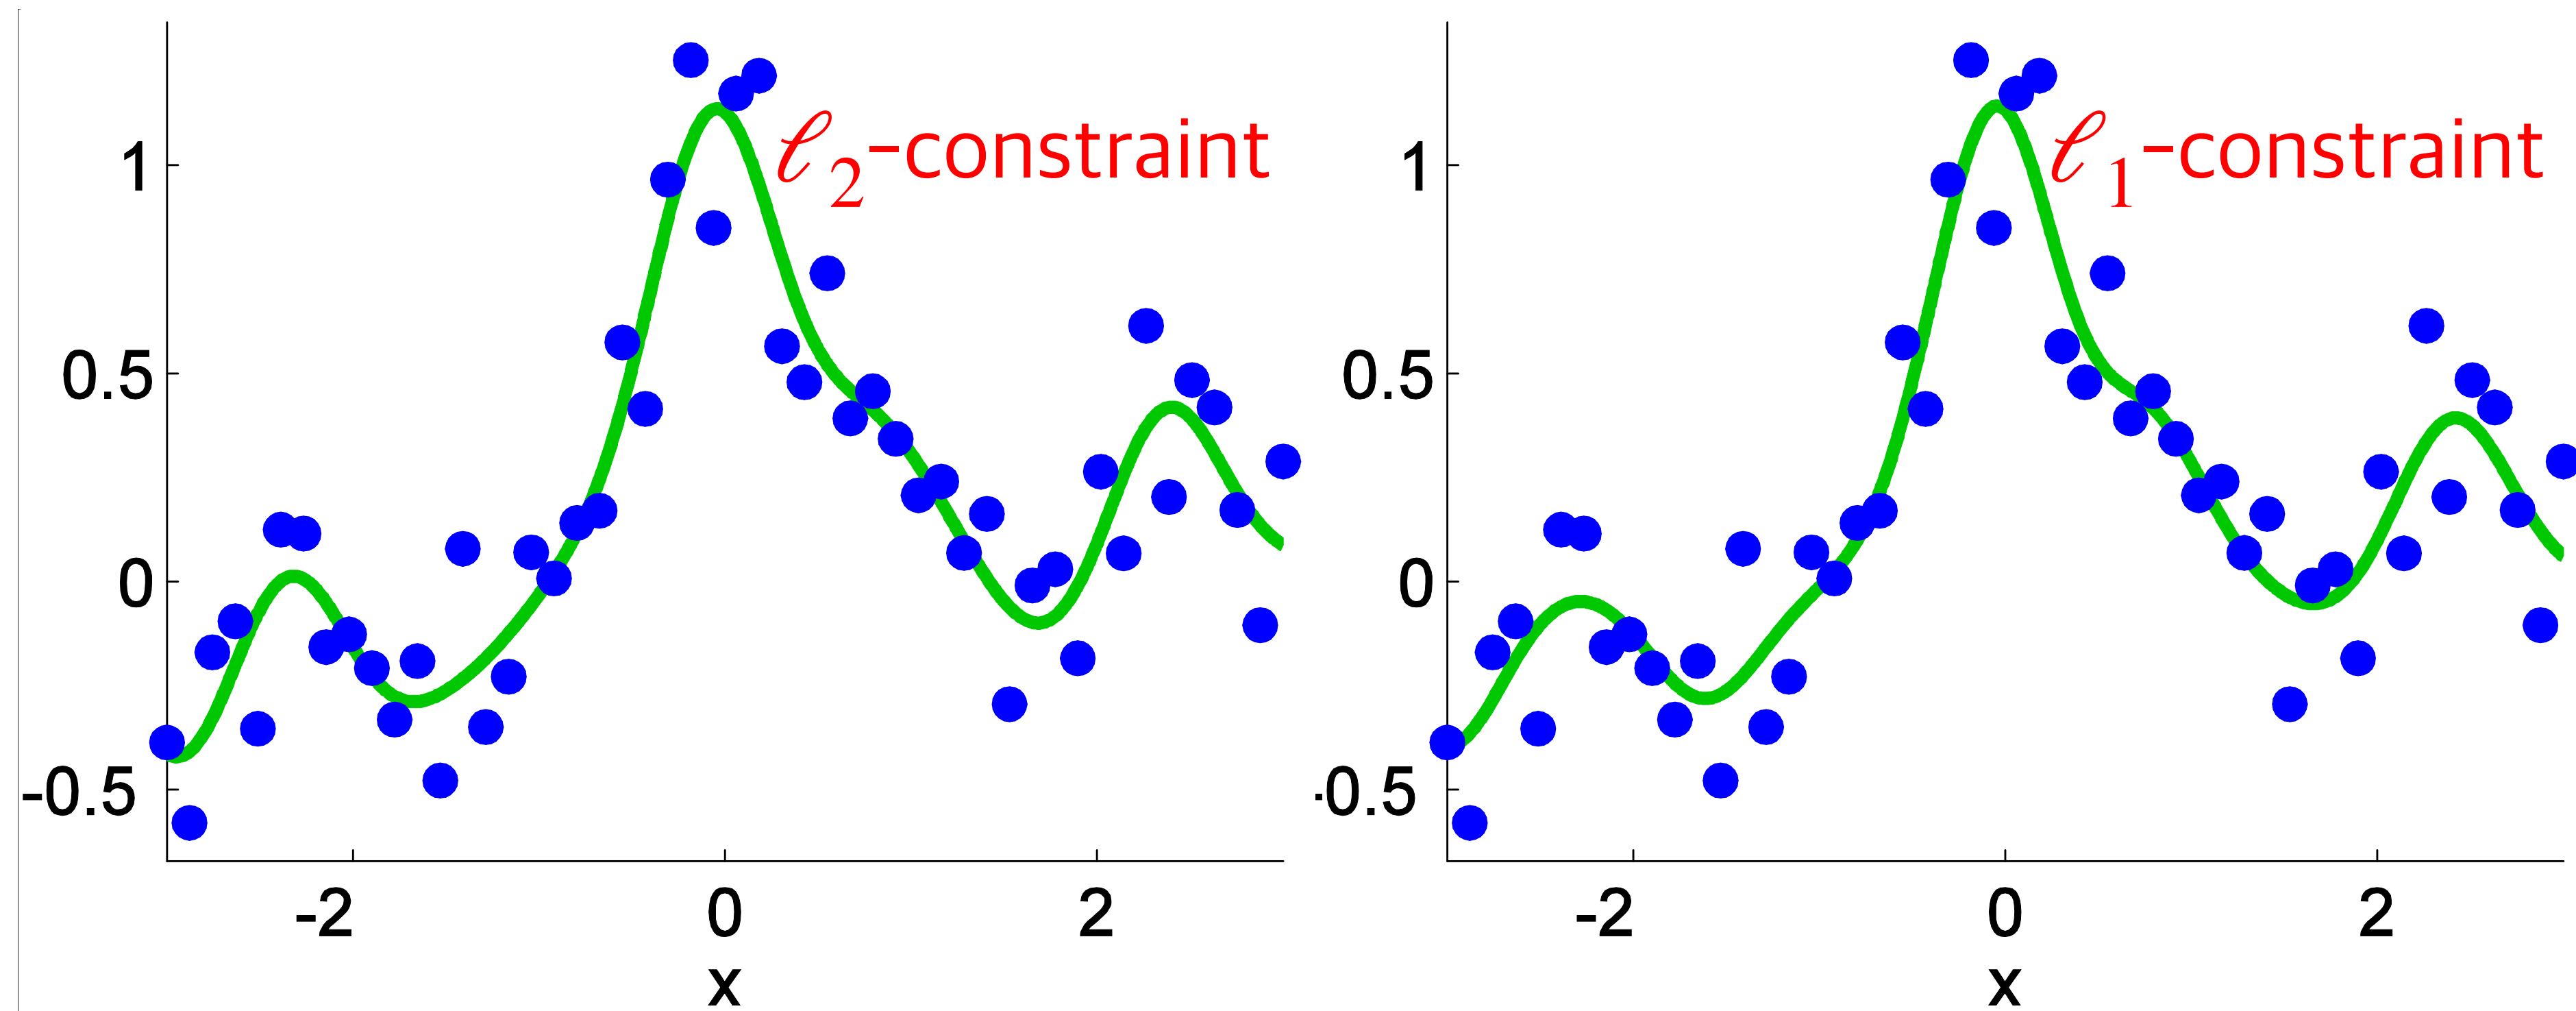
$$J(\tilde{\theta}) = \tilde{J}(\tilde{\theta}) \geq \tilde{J}(\hat{\theta}) \geq J(\hat{\theta})$$



# Example

$$f_{\theta}(\mathbf{x}) = \sum_{j=1}^n \theta_j \exp \left( -\frac{\|\mathbf{x} - \mathbf{x}_j\|^2}{2h^2} \right)$$

Implementing  
this is  
homework!




- $\ell_1$ -constraint results look roughly the same as  $\ell_2$ -constraint
- However, the  $\ell_1$ -constraint has 37 out of 50 parameters that are zero!

Note: if absolute value is below  $1e-3$ , we regard it as zero.

# Issues of iteratively reweighted shrinkage

- The cost of 1 iteration is heavy and when  $c_j$  becomes extremely small, may become unstable due to  $1/c_j$ .
- Still, the super simple implementation makes it a practical choice.
- Many advanced methods exists, including:
  - Accelerated proximal gradient method
  - Alternating direction method of multipliers (ADMM)



# Contents

26

## 1. Sparse regression

1.  $\ell_1$ -constrained least squares regression

2. Solving  $\ell_1$ -constrained LS

3. Various extensions

A) Online learning

B) Generalized  $\ell_1$ -norm

C)  $\ell_p$ -norm

D)  $\ell_1 + \ell_2$ -norm

E)  $\ell_{1,2}$ -norm

## 2. Robust regression



$$\min_{\boldsymbol{\theta}} \frac{1}{2} \sum_{i=1}^n \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2$$

$$f_{\boldsymbol{\theta}}(\mathbf{x}) = \sum_{j=1}^b \theta_j \phi_j(\mathbf{x})$$

1. Initialize parameter  $\boldsymbol{\theta}$ .
2. Gradient descent on a randomly chosen  $(\mathbf{x}, y)$ :

$$\boldsymbol{\theta} \longleftarrow \boldsymbol{\theta} - \varepsilon \frac{\partial}{\partial \boldsymbol{\theta}} \frac{(f_{\boldsymbol{\theta}}(\mathbf{x}) - y)^2}{2}$$

$$= \boldsymbol{\theta} - \varepsilon \boldsymbol{\phi}(\mathbf{x}) \left( \boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x}) - y \right)$$

stepsize  
 $\varepsilon > 0$

3. Repeat step 2. until convergence
- How can we satisfy the  $\ell_1$  constraint?

# Sparse online learning

28

$$\min_{\boldsymbol{\theta}} \frac{1}{2} \sum_{i=1}^n \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2 + \lambda \sum_{j=1}^b |\theta_j| \quad f_{\boldsymbol{\theta}}(\mathbf{x}) = \sum_{j=1}^b \theta_j \phi_j(\mathbf{x})$$

1. Initialize parameter  $\boldsymbol{\theta}$ .
2. Gradient descent on a randomly chosen  $(\mathbf{x}, y)$ :

$$\boldsymbol{\theta} \longleftarrow \boldsymbol{\theta} - \varepsilon \phi(\mathbf{x}) \left( \boldsymbol{\theta}^\top \phi(\mathbf{x}) - y \right) \quad \begin{array}{l} \text{stepsize} \\ \varepsilon > 0 \end{array}$$

3. Make the solution sparse:


$$\forall j = 1, \dots, b, \quad \theta_j \longleftarrow \begin{cases} \max(0, \theta_j - \lambda \varepsilon) & (\theta_j > 0), \\ \min(0, \theta_j + \lambda \varepsilon) & (\theta_j \leq 0). \end{cases}$$

4. Repeat steps 2, 3 until convergence.


# Intuition of making it sparse

$$\theta_j \leftarrow \begin{cases} \max(0, \theta_j - \lambda\varepsilon) & (\theta_j > 0), \\ \min(0, \theta_j + \lambda\varepsilon) & (\theta_j \leq 0). \end{cases}$$

$$\max(0, \theta_j - \lambda\varepsilon) + \min(0, \theta_j + \lambda\varepsilon)$$



- With regularization, we are closer to the origin.
- The method of performing corrections corresponding to the regularization term for stochastic gradient descent is called proximal gradient method.



# Contents

30

1. Sparse regression
  1.  $\ell_1$ -constrained least squares regression
  2. Solving  $\ell_1$ -constrained LS
  3. Various extensions
    - A) Online learning
    - B) Generalized  $\ell_1$ -norm
    - C)  $\ell_p$ -norm
    - D)  $\ell_1 + \ell_2$ -norm
    - E)  $\ell_{1,2}$ -norm
2. Robust regression

# Generalized $\ell_1$ -norm

$$\|\mathbf{F}\boldsymbol{\theta}\|_1 = \sum_j \left| \sum_{j'} F_{j,j'} \theta_{j'} \right|$$

- **Example:** the norm of the difference of adjacent elements

$$\sum_j |\theta_{j+1} - \theta_j| \quad F_{j,j'} = \begin{cases} 1 & (j' = j + 1) \\ -1 & (j' = j) \\ 0 & (\text{otherwise}) \end{cases}$$

- The generalized  $\ell_1$  version of LS is called **fused lasso**.

# Generalized $\ell_1$ -constrained least squares regression 32


$$\min_{\boldsymbol{\theta}} \frac{1}{2} \sum_{i=1}^n \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2 \quad \text{subject to } \|\mathbf{F}\boldsymbol{\theta}\|_1 \leq R \quad R \geq 0$$

■ **Example:** total variation noise removal

$$\min_{\boldsymbol{\theta}} \frac{1}{2} \|\boldsymbol{\theta} - \mathbf{y}\|^2 \quad \text{subject to } \sum_j |\theta_{j+1} - \theta_j| \leq R$$

From Wikipedia

$$F_{j,j'} = \begin{cases} 1 & (j' = j + 1) \\ -1 & (j' = j) \\ 0 & (\text{otherwise}) \end{cases}$$







# Contents

33

1. Sparse regression
  1.  $\ell_1$ -constrained least squares regression
  2. Solving  $\ell_1$ -constrained LS
  3. Various extensions
    - A) Online learning
    - B) Generalized  $\ell_1$ -norm
    - C)  $\ell_p$ -norm
    - D)  $\ell_1 + \ell_2$ -norm
    - E)  $\ell_{1,2}$ -norm
2. Robust regression


$\ell_p$ -norm

$$\|\boldsymbol{\theta}\|_p = \left( \sum_{j=1}^b |\theta_j|^p \right)^{\frac{1}{p}}$$

$$p \geq 1$$

$$\|\boldsymbol{\theta}\|_p = \sum_{j=1}^b |\theta_j|^p$$

$$p \leq 1$$



$$p = 0$$


$$\|\boldsymbol{\theta}\|_0 = \# \text{ non-zero elements}$$

$$p = \infty$$

$$\|\boldsymbol{\theta}\|_\infty = \max \{ |\theta_1|, \dots, |\theta_b| \}$$

- Restrict the model to be within  $\ell_p$ -hypercube


$$\min_{\boldsymbol{\theta}} \frac{1}{2} \sum_{i=1}^n \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2 \quad \text{subject to } \|\boldsymbol{\theta}\|_p \leq R \quad p \geq 0$$




# $\ell_p$ -constrained least squares regression

36

- The solution becomes sparse:  $0 \leq p \leq 1$
- Optimization problem is convex:  $p \geq 1$   
(Easier to achieve global solution)
- Only  $p = 1$  satisfies both!





# Contents

37

## 1. Sparse regression

1.  $\ell_1$ -constrained least squares regression

2. Solving  $\ell_1$ -constrained LS

3. Various extensions

A) Online learning

B) Generalized  $\ell_1$ -norm

C)  $\ell_p$ -norm

D)  $\ell_1 + \ell_2$ -norm

E)  $\ell_{1,2}$ -norm

## 2. Robust regression

## Issue of $\ell_1$ -constrained LS regression

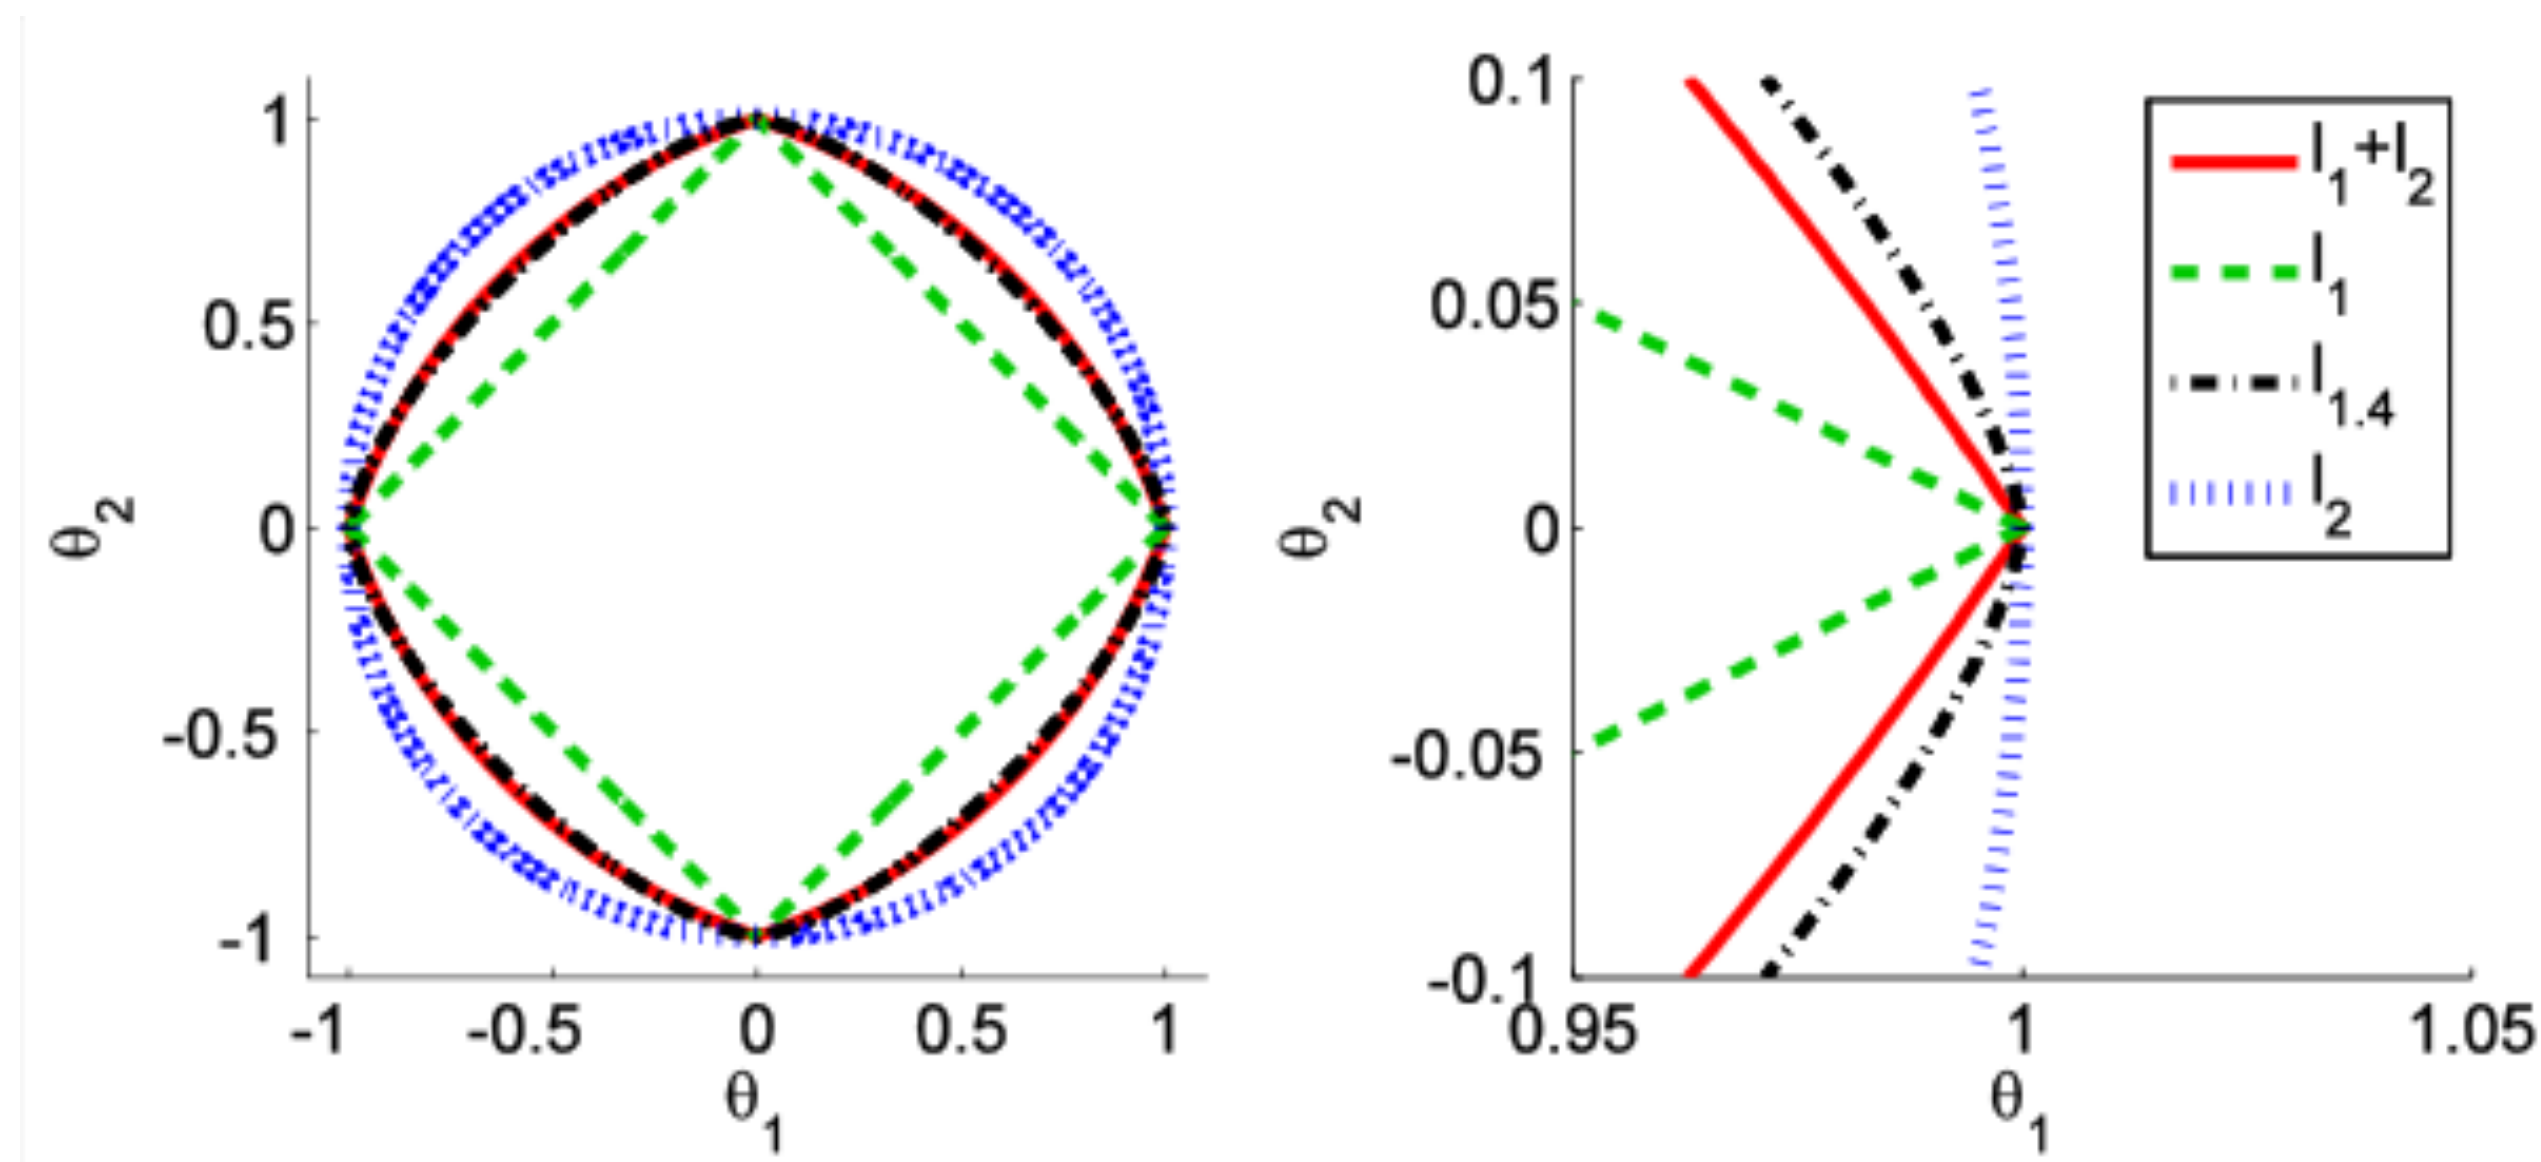
$$\min_{\boldsymbol{\theta}} \frac{1}{2} \sum_{i=1}^n \left( \sum_{j=1}^b \theta_j \phi_j(\mathbf{x}_i) - y_i \right)^2 \quad \text{subject to} \quad \sum_{j=1}^b |\theta_j| \leq R$$

- When some basis functions  $\{\phi_j(\mathbf{x})\}_{j=1}^b$  are similar, only one of them is chosen.
- When  $b < n$ , this may perform worse compared with  $\ell_2$ -constrained LS regression.

# $\ell_1 + \ell_2$ -constrained least squares regression 39

$$(1 - \tau) \sum_{j=1}^b |\theta_j| + \tau \sum_{j=1}^b \theta_j^2 \leq R \quad 0 \leq \tau < 1$$

- Similar to  $\ell_{1.4}$ -cube, but  $\ell_1 + \ell_2$ -cube is sharp



- Also called **elastic net**.





# Contents

40

## 1. Sparse regression

1.  $\ell_1$ -constrained least squares regression

2. Solving  $\ell_1$ -constrained LS

3. Various extensions

A) Online learning

B) Generalized  $\ell_1$ -norm

C)  $\ell_p$ -norm

D)  $\ell_1 + \ell_2$ -norm

E)  $\ell_{1,2}$ -norm

## 2. Robust regression

## $\ell_{1,2}$ -norm



- Setup: there are variable groups that are similar, but we do not know whether they are helpful.
- When  $\boldsymbol{\theta} = (\theta_1, \dots, \theta_b)^\top$  has a group structure  $\boldsymbol{\theta} = (\boldsymbol{\theta}^{(1)\top}, \dots, \boldsymbol{\theta}^{(t)\top})^\top$ , then the following is called  $\ell_{1,2}$ -norm:

$$\|\boldsymbol{\theta}\|_{1,2} = \sum_{j=1}^t \|\boldsymbol{\theta}^{(j)}\|_2 \quad \boldsymbol{\theta}^{(j)} \in \mathbb{R}^{b_j}$$

- Called **group regularization** when this is used for the regularization term.

# Exercise


- The constraint of  $\ell_2$ -norm and  $\ell_1$ -norm in 3-dimensional case:



- Similarly, visualize the constraint of  $\ell_{1,2}$ -norm and the sparse solution:

$$\sqrt{\theta_1^2 + \theta_2^2} + \sqrt{\theta_3^2} \leq R$$

# Visualization



- Achieve sparsity at the group level!

$$\sum_{j=1}^t \|\boldsymbol{\theta}^{(j)}\| \leq R$$

# Summary of sparse regression


44

- Achieve sparsity by  $\ell_1$ -regularized learning.
- We no longer have an analytical solution.
- We can consider many different ideas: fused lasso, group lasso, elastic net, ...
- Sparse regression is used widely in many domains.


# Another Issue of Least Squares

- With just a single outlier, the results are altered heavily:

$$f_{\theta}(x) = \theta_1 + \theta_2 x$$



LS regression  
(w/o outlier)




LS regression  
(w/ outlier)




# $\ell_2$ loss function

$$\sum_{i=1}^n \left( f_{\theta}(\mathbf{x}_i) - y_i \right)^2$$



- **Least squares regression**: measures the goodness of fit of the training output by the  $\ell_2$  loss function.
- The influence of outliers is strengthened by the “squared” component.
- Need to reduce the influence of outliers to stabilize learning!



# Contents


47

1. Sparse regression
2. Robust regression
  1.  $\ell_1$ -loss
    1. Relationship between median
    2. Robustness and estimation accuracy
  2. Huber loss
  3. Tukey loss





# Robust regression with $\ell_1$ loss

$$\min_{\theta} \sum_{i=1}^n |f_{\theta}(x_i) - y_i|$$


- Goodness of fit to training samples is measured by  $\ell_1$  function:
- The influence of the outlier becomes linear.
- Called least absolute (LA).



# Example



Will explain  
the  
methods  
soon!



# Contents

50

1. Sparse regression
2. Robust regression
  1.  $\ell_1$ -loss
    1. Relationship between median
    2. Robustness and estimation accuracy
  2. Huber loss
  3. Tukey loss

- Probability of a continuous random variable taking a value less than or equal to  $x$

$$P(x) = \text{Prob}(X \leq x) = \int_{-\infty}^x p(u) du$$

$P(x)$ : cumulative distribution function

$$P'(x) = \frac{dP(x)}{dx} = p(x)$$

- (Weakly) increasing function:

$$x_1 < x_2 \implies P(x_1) \leq P(x_2)$$

- Range:

$$x \rightarrow -\infty \implies P(x) \rightarrow 0$$

$$x \rightarrow \infty \implies P(x) \rightarrow 1$$





# Expectation and Median

## ■ Expectation:


$$E[X] = \int xp(x)dx$$

## ■ Median:

$$x \text{ satisfying } \text{Prob}(X \leq x) = \frac{1}{2}$$



- Expected values can be inconsistent with our intuition when there are outliers.



Example: With one wealthy person, everyone else falls below the expectation. (Sometimes referred as "Bill Gates walks into a bar" scenario.)

↓ Single outlier


# Math exercise

- Consider density function  $p(y)$  with support  $[a, b]$ .
- Show  $\theta_2$  that minimizes **squared error**  $J_2(y)$  is the **expectation** of  $y$ .

$$\theta_2 = \operatorname{argmin}_{\theta} J_2(\theta)$$

$$J_2(\theta) = \int_a^b (y - \theta)^2 p(y) dy$$

- **Helpful info:** derivative is zero at the minimum.



# Math exercise

- Consider density function  $p(y)$  with support  $[a, b]$ .
- $\theta_1$  that minimizes the absolute error  $J_1(\theta)$  is the median

$$\theta_1 = \operatorname{argmin}_{\theta} J_1(\theta) \quad J_1(\theta) = \int_a^b |y - \theta| p(y) dy$$

- Helpful info: integration by parts

$$\int_a^b f(y) g'(y) dy = \left[ f(y) g(y) \right]_a^b - \int_a^b f'(y) g(y) dy$$

- Observed value = true value + noise:


$$\{y_i \mid y_i = \mu^* + \epsilon_i\}_{i=1}^n$$

- $\ell_2$ -loss function: corresponds to the estimation of the **expectation** of the observations.

- $$\operatorname{argmin}_{\mu} \left[ \sum_{i=1}^n (y_i - \mu)^2 \right] = \operatorname{mean} \left( \{y_i\}_{i=1}^n \right)$$

- $\ell_1$ -loss function: corresponds to the estimation of **median** of the observations.

- $$\operatorname{argmin}_{\mu} \left[ \sum_{i=1}^n |y_i - \mu| \right] = \operatorname{median} \left( \{y_i\}_{i=1}^n \right)$$




# Contents

56


1. Sparse regression
2. Robust regression
  1.  $\ell_1$ -loss
    1. Relationship between median
    2. Robustness and estimation accuracy
  2. Huber loss
  3. Tukey loss

- **Breakpoint**: the proportion of samples that will maintain non-breaking solutions when we replace samples with infinity.
  - $\ell_2$ -loss function: 0%  
(💣 not robust)
  - $\ell_1$ -loss function: 50%  
(😎 robust)
- However, the  $\ell_1$ -loss function is **not efficient** for Gaussian noise (large variance).





- Highly robust = ignoring more training samples
  - A regressor that always outputs 0 is meaningless but super robust.
- Practical requirements:
  - Want to be similar to LS when there are no outliers.
  - Need robustness when there are many or large outliers.



# Contents


59

1. Sparse regression
2. Robust regression
  1.  $\ell_1$ -loss
  2. Huber loss
  3. Tukey loss

# Huber loss

- The best of both world?
  - Squared for small errors
  - Absolute for larger errors

$$\min_{\theta} \sum_{i=1}^n \rho(f_{\theta}(\mathbf{x}_i) - y_i)$$




$$\rho(r) = \begin{cases} r^2/2 & (|r| \leq \eta) \\ \eta|r| - \eta^2/2 & (|r| > \eta) \end{cases} \quad \eta \geq 0$$


- Parameter  $\eta$  is designed by the user as the point at which errors may originate from outliers.

# 1st Approach

- Huber loss is continuously differentiable.



$$\rho(r) = \begin{cases} r^2/2 & (|r| \leq \eta) \\ \eta|r| - \eta^2/2 & (|r| > \eta) \end{cases}$$




$$\rho'(r) = \begin{cases} r & (|r| \leq \eta) \\ \text{sign}(r)\eta & (|r| > \eta) \end{cases}$$

- Gradient method:

$$\theta \leftarrow \theta - \varepsilon \nabla J(\theta)$$


$$J(\theta) = \sum_{i=1}^n \rho(f_{\theta}(x_i) - y_i)$$



# 2nd Approach

62

- Gradient methods are tricky to adjust step widths
- Iterative least squares algorithm:
  - Huber loss is suppressed from above by a quadratic function tangent to the current solution (different from Newton's method)
  - Minimizing the quadratic upper bound analytically to find a better solution step by step




# Math exercise

- Huber loss when  $|r| > \eta$ :  $\eta|r| - \frac{\eta^2}{2}$
- $$\rho(r) = \begin{cases} r^2/2 & (|r| \leq \eta) \\ \eta|r| - \eta^2/2 & (|r| > \eta) \end{cases}$$

Derive a quadratic function that is tangent to the huber loss.

- Helpful info: from symmetry, quadratic function that is tangent at  $\pm c$  can be expressed as  $ar^2 + b$



# Minimization of upper bound

$$\rho(r) = \begin{cases} r^2/2 & (|r| \leq \eta) \\ \eta|r| - \eta^2/2 & (|r| > \eta) \end{cases}$$

- Upper bound  $\tilde{\rho}(r) \geq \rho(r)$  for residual  $\tilde{r} = f_{\tilde{\theta}}(\mathbf{x}) - y$  (where  $\tilde{\theta}$  is current param):

$$\tilde{\rho}(r) = \begin{cases} r^2/2 & (|\tilde{r}| \leq \eta) \\ \frac{\eta}{2|\tilde{r}|} r^2 + \underbrace{\frac{\eta|\tilde{r}|}{2} - \frac{\eta^2}{2}}_{\text{constant}} & (|\tilde{r}| > \eta) \end{cases}$$

$$= \frac{\tilde{w}}{2} r^2 + \text{const}$$

$$\tilde{w} = \begin{cases} 1 & (|\tilde{r}| \leq \eta) \\ \eta/|\tilde{r}| & (|\tilde{r}| > \eta) \end{cases}$$



- Original objective we wanted to minimize:

$$J(\boldsymbol{\theta}) = \sum_{i=1}^n \rho(f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i) \quad \rho(r) = \begin{cases} r^2/2 & (|r| \leq \eta) \\ \eta|r| - \eta^2/2 & (|r| > \eta) \end{cases}$$


- Minimization of  $\tilde{J}$  (upper bound of  $J$ ) derived with  $\tilde{\boldsymbol{\theta}}$ :

$$\hat{\boldsymbol{\theta}} = \underset{\boldsymbol{\theta}}{\operatorname{argmin}} \tilde{J}(\boldsymbol{\theta}) \quad \tilde{J}(\boldsymbol{\theta}) = \frac{1}{2} \sum_{i=1}^n \tilde{w}_i \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2$$

$$\tilde{w}_i = \begin{cases} 1 & (|\tilde{r}_i| \leq \eta) \\ \eta/|\tilde{r}_i| & (|\tilde{r}_i| > \eta) \end{cases} \quad \tilde{r}_i = f_{\tilde{\boldsymbol{\theta}}}(\mathbf{x}_i) - y_i$$

- Upper bound is weighted least squares:

$$\min_{\theta} \frac{1}{2} \sum_{i=1}^n \tilde{w}_i \left( f_{\theta}(\mathbf{x}_i) - y_i \right)^2$$



- Small weights are set for outliers

$$\tilde{w}_i = \begin{cases} 1 & (|\tilde{r}_i| \leq \eta) \\ \eta/|\tilde{r}_i| & (|\tilde{r}_i| > \eta) \end{cases} \quad \tilde{r}_i = f_{\tilde{\theta}}(\mathbf{x}_i) - y_i$$


- The minimum solution of the upper bound is analytically obtained by  $\hat{\theta} = (\Phi^{\top} \tilde{\mathbf{W}} \Phi)^{-1} \Phi^{\top} \tilde{\mathbf{W}} \mathbf{y}$

$$\tilde{\mathbf{W}} = \text{diag}(\tilde{w}_1, \dots, \tilde{w}_n)$$

Proving this is homework.

- Since the upper bound is tangent at  $\tilde{\theta}$ :  $J(\tilde{\theta}) = \tilde{J}(\tilde{\theta})$
- If  $\hat{\theta}$  is the minimal solution of the upper bound:  $\tilde{J}(\tilde{\theta}) \geq \tilde{J}(\hat{\theta})$
- Since  $\tilde{J}$  is the upper bound of  $J$ :  $\tilde{J}(\hat{\theta}) \geq J(\hat{\theta})$
- To summarize:  $J(\tilde{\theta}) = \tilde{J}(\tilde{\theta}) \geq \tilde{J}(\hat{\theta}) \geq J(\hat{\theta})$

💡 when updating from  $\tilde{\theta}$  to  $\hat{\theta}$ ,  $J$  generally decreases



$$\hat{\theta} = \operatorname{argmin}_{\theta} \tilde{J}(\theta)$$

- Initialize  $\theta$ .
- Repeat below until convergence:
  - Derive  $\mathbf{W}$  from current solution  $\theta$  (derive upper bound)

$$\mathbf{W} = \text{diag}(w_1, \dots, w_n)$$

$$w_i = \begin{cases} 1 & (|f_{\theta}(\mathbf{x}_i) - y_i| \leq \eta) \\ \eta / |f_{\theta}(\mathbf{x}_i) - y_i| & (|f_{\theta}(\mathbf{x}_i) - y_i| > \eta) \end{cases}$$

- Update  $\theta$  (minimize upper bound)

$$\theta \leftarrow (\Phi^{\top} \mathbf{W} \Phi)^{-1} \Phi^{\top} \mathbf{W} \mathbf{y}$$

- Iterative least squares algorithm for Huber regression for the linear model  $f_{\theta}(x) = \theta_1 + \theta_2 x$ .

```
import numpy as np
import matplotlib

matplotlib.use('TkAgg')
import matplotlib.pyplot as plt

np.random.seed(1)

def generate_sample(x_min=-3., x_max=3., sample_size=10):
    x = np.linspace(x_min, x_max, num=sample_size)
    y = x + np.random.normal(loc=0., scale=.2, size=sample_size)
    y[-1] = -4 # outlier
    return x, y

def build_design_matrix(x):
    phi = np.empty(x.shape + (2,))
    phi[:, 0] = 1.
    phi[:, 1] = x
    return phi


def iterative_reweighted_least_squares(x, y, eta=1., n_iter=1000):
    phi = build_design_matrix(x)
    # initialize theta using the solution of regularized ridge regression
    theta = theta_prev = np.linalg.solve(
        phi.T.dot(phi) + 1e-4 * np.identity(phi.shape[1]), phi.T.dot(y))
    for _ in range(n_iter):
        r = np.abs(phi.dot(theta_prev) - y)
        w = np.diag(np.where(r > eta, eta / r, 1.))
        phit_w_phi = phi.T.dot(w).dot(phi)
        phit_w_y = phi.T.dot(w).dot(y)
        theta = np.linalg.solve(phit_w_phi, phit_w_y)
        if np.linalg.norm(theta - theta_prev) < 1e-3:
            break
        theta_prev = theta
    return theta
```



# Python implementation 2

70

```
def predict(x, theta):  
    phi = build_design_matrix(x)  
    return phi.dot(theta)  
  
def visualize(x, y, theta, x_min=-4., x_max=4., filename='xxxxxxx.png'):  
    X = np.linspace(x_min, x_max, 1000)  
    Y = predict(X, theta)  
    plt.clf()  
    plt.plot(X, Y, color='green')  
    plt.scatter(x, y, c='blue', marker='o')  
    plt.savefig(filename)  
  
x, y = generate_sample()  
theta = iterative_reweighted_least_squares(x, y, eta=1.)  
visualize(x, y, theta)
```



# Contents

71

1. Sparse regression
2. Robust regression
  1.  $\ell_1$ -loss
  2. Huber loss
  3. Tukey loss




# Can we improve further?

- Huber loss is robust compared with  $\ell_2$ -loss
- But... no upper bound on the loss, so the effect of outliers remain to some extent.

$$\min_{\theta} \sum_{i=1}^n \rho_{\text{Huber}}(f_{\theta}(\mathbf{x}_i) - y_i)$$

$$\rho_{\text{Huber}}(r) = \begin{cases} r^2/2 & (|r| \leq \eta) \\ \eta|r| - \eta^2/2 & (|r| > \eta) \end{cases}$$




# Tukey loss

$$r = f_{\theta}(x) - y$$

- Consider an upper bounded loss

$$\rho_{\text{Tukey}}(r) = \begin{cases} \left(1 - [1 - r^2/\eta^2]^3\right) / 6 & (|r| \leq \eta) \\ 1/6 & (|r| > \eta) \end{cases}$$



$$\min_{\theta} \sum_{i=1}^n \rho(f_{\theta}(\mathbf{x}_i) - y_i)$$

- Consider differentiable and symmetric loss  $\rho(r)$ .
- Quadratic upper bound that is tangent to  $\rho(r)$  at  $\pm \tilde{r}$ :

$$\tilde{\rho}(r) = \frac{\tilde{w}}{2} r^2 + \text{const}$$

$$\tilde{w} = \rho'(\tilde{r})/\tilde{r}$$

(You can try to prove this yourself)

- Iterative LS algorithm:

$$\min_{\theta} \frac{1}{2} \sum_{i=1}^n \tilde{w}_i \left( f_{\theta}(\mathbf{x}_i) - y_i \right)^2$$

$$\tilde{w}_i = \rho'(\tilde{r}_i)/\tilde{r}_i$$

$$\tilde{r}_i = f_{\tilde{\theta}}(\mathbf{x}_i) - y_i$$

# Weight function for Tukey loss

- Tukey loss:

$$\rho_{\text{Tukey}}(r) = \begin{cases} \left(1 - [1 - r^2/\eta^2]^3\right) / 6 & (|r| \leq \eta) \\ 1/6 & (|r| > \eta) \end{cases}$$


- Weight function for Tukey loss:

$$\tilde{w} = \rho'(\tilde{r})/\tilde{r}$$

$$w_{\text{Tukey}} = \begin{cases} (1 - r^2/\eta^2)^2 & (|r| \leq \eta) \\ 0 & (|r| > \eta) \end{cases}$$



- Weight is zero for large residuals

$$w_{\text{Huber}} = \begin{cases} 1 & (|r| \leq \eta) \\ \eta/|r| & (|r| > \eta) \end{cases}$$



# Simple example


- Tukey regression is more robust to outliers.
- Since it is a non-convex optimization problem, the solution may depend on initialization.



# Summary of robust regression

77

- Squared loss (mean) is vulnerable to outliers
- Absolute loss (median) is robust to outliers
- Huber loss balances robustness and efficiency
  - Solution cannot be obtained analytically
- Tukey loss improves robustness further
  - May need to be a bit more cautious about optimization



# Summary of regression

78

1. Learning models
2. Least squares regression
3. Regularized regression
4. Sparse regression
5. Robust regression



# Summary of regression

79

- Models for learning functions
  - Linear models
  - Kernel models
  - Non-linear models
- Least-squares regression
  - Minimizes the squared error with the training sample
  - Solutions can be calculated analytically
- Online regression
  - Sequential learning by extracting data one at a time
  - Large amounts of data can be handled

# Summary of regression


80

- $\ell_2$ -regularized regression
  - Reduces over-fitting of least-squares regression
  - Solutions can be calculated analytically
  - Cross-validation is used for model selection
- $\ell_1$ -regularized regression (sparse regression)
  - Data where many of the true parameter values are zero can be trained properly.
- Robust regression
  - Enhanced robustness against outliers.

# Schedule

81

- |    |       |                  |     |       |                           |
|----|-------|------------------|-----|-------|---------------------------|
| 1. | 04/8  | Introduction     | 8.  | 06/17 | Deep learning 3           |
| 2. | 04/15 | Regression 1     | 9.  | 06/24 | Semi-supervised learning  |
| 3. | 04/22 | Regression 2     |     |       |                           |
| ●  | 04/30 | Cancelled        | 10. | 07/01 | Language models           |
| 4. | 05/13 | Classification 1 | 11. | 07/08 | Representation learning 1 |
| 5. | 05/20 | Classification 2 |     |       |                           |
| 6. | 05/27 | Deep learning 1  | 12. | 07/15 | Representation learning 2 |
| ●  | 06/03 | No lecture       |     |       |                           |
| 7. | 06/10 | Deep learning 2  | 13. | 07/22 | Advanced topics           |



# Coming up next

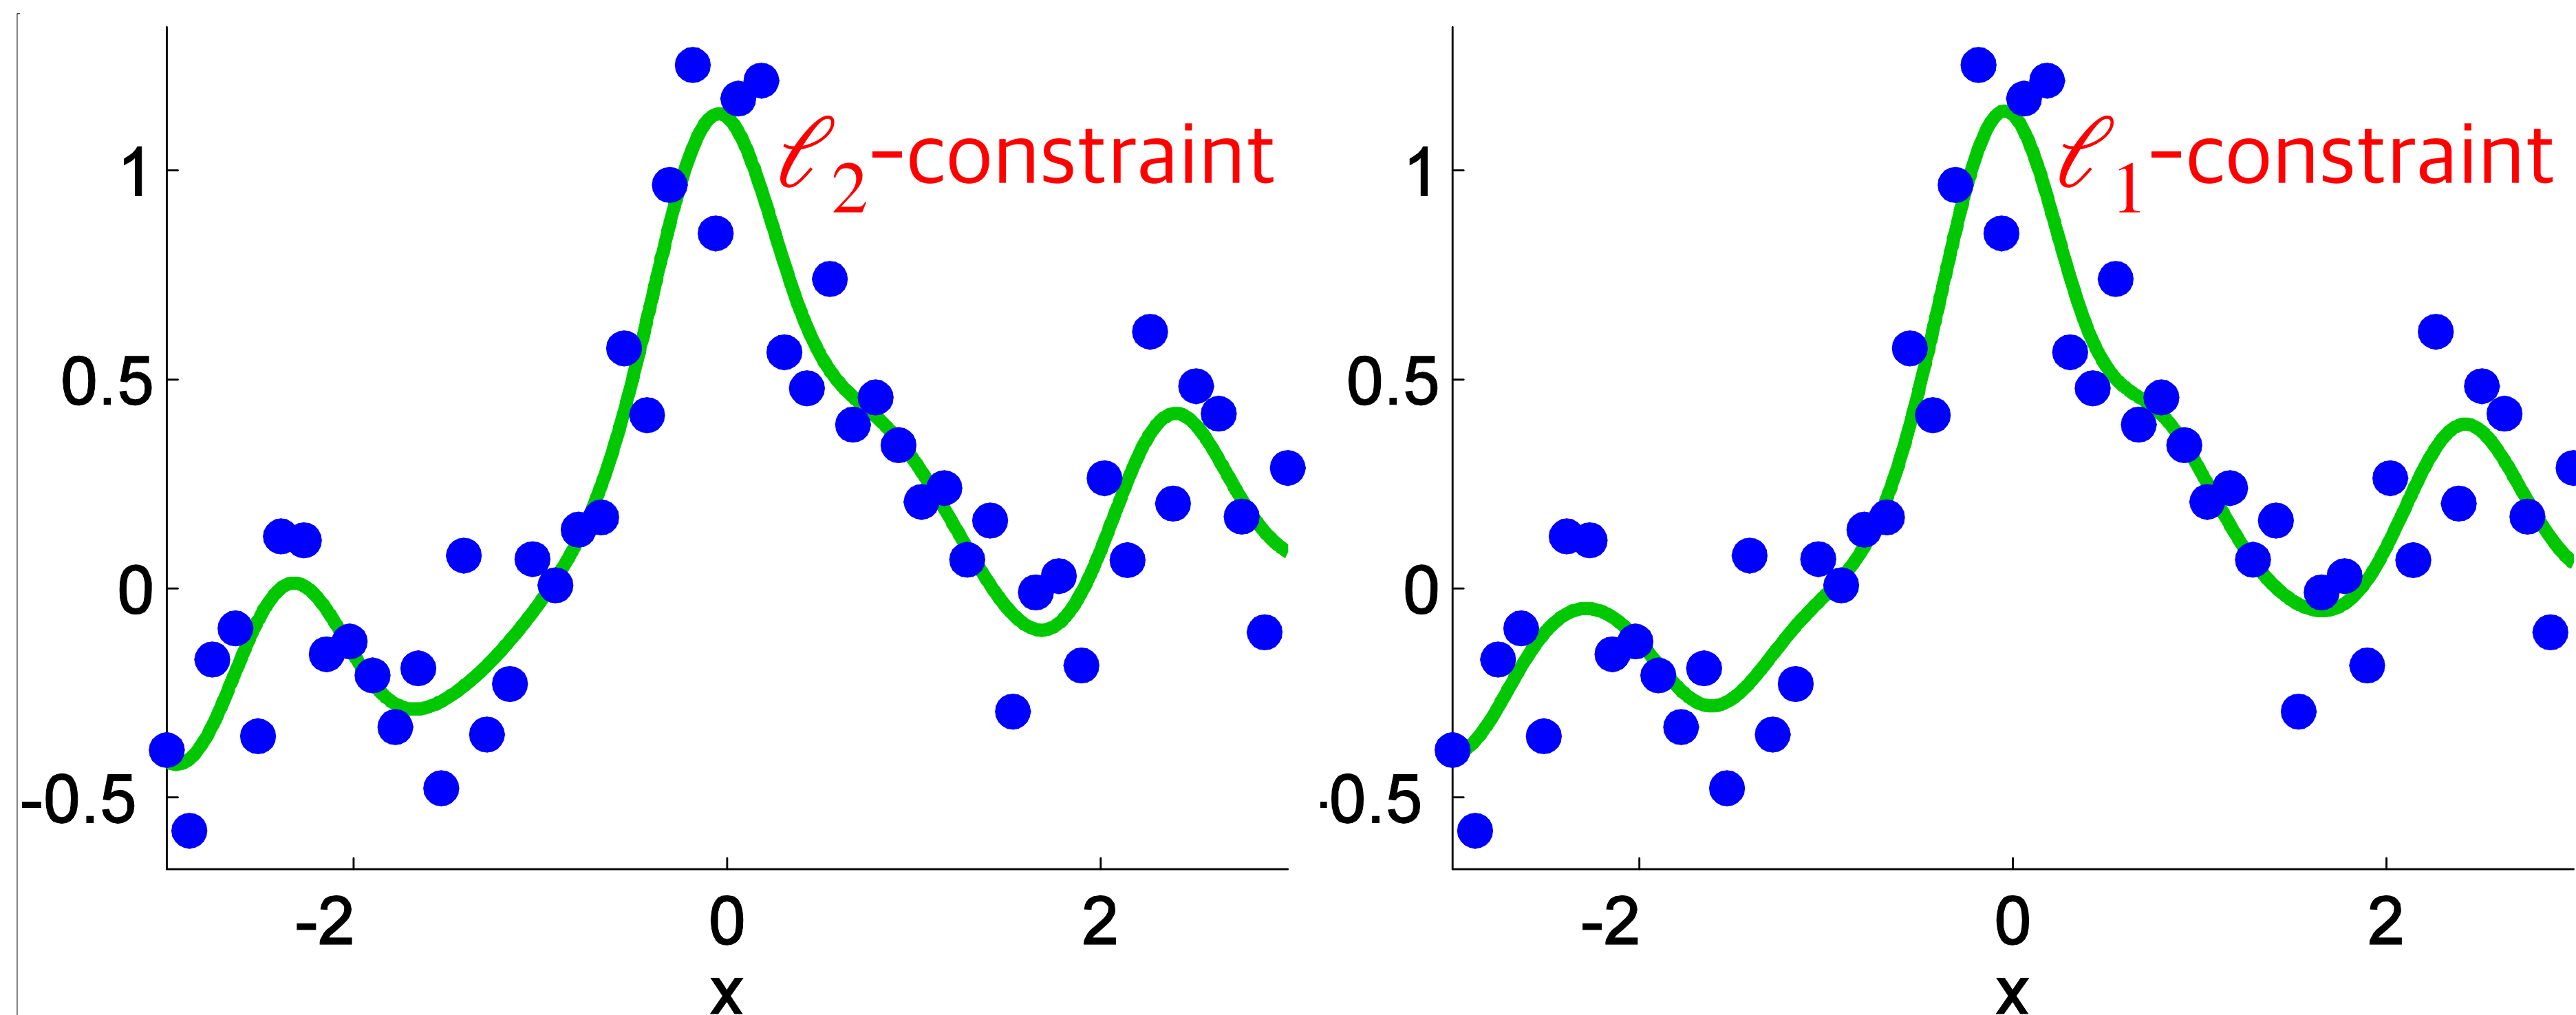
82

- Classification 1

# Homework 1

- Implement the  $\ell_1$ -constraint LS (iteratively reweighted shrinkage). You may use the same data from previous week. (No need to do cross validation).

$$f_{\theta}(\mathbf{x}) = \sum_{j=1}^n \theta_j \exp \left( -\frac{\|\mathbf{x} - \mathbf{x}_j\|^2}{2h^2} \right)$$



- Setup

- Linear model:  $f_{\boldsymbol{\theta}}(\mathbf{x}) = \sum_{j=1}^b \theta_j \phi_j(\mathbf{x})$

- Basis functions:  $\{\phi_j(\mathbf{x})\}_{j=1}^b$

- Prove that the solution to the weighted LS method is the following:  $\hat{\boldsymbol{\theta}} = (\boldsymbol{\Phi}^\top \widetilde{\mathbf{W}} \boldsymbol{\Phi})^{-1} \boldsymbol{\Phi}^\top \widetilde{\mathbf{W}} \mathbf{y}$

- Weighted LS problem:  $\min_{\boldsymbol{\theta}} \frac{1}{2} \sum_{i=1}^n \tilde{w}_i \left( f_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)^2$

$$\boldsymbol{\Phi} = \begin{pmatrix} \phi_1(\mathbf{x}_1) & \cdots & \phi_b(\mathbf{x}_1) \\ \vdots & \ddots & \vdots \\ \phi_1(\mathbf{x}_n) & \cdots & \phi_b(\mathbf{x}_n) \end{pmatrix}$$

$$\widetilde{\mathbf{W}} = \text{diag}(\tilde{w}_1, \dots, \tilde{w}_n)$$

$$\mathbf{y} = (y_1, \dots, y_n)^\top$$





